

# ActivePointers: A Case for Software Address Translation on GPUs

*Sagi Shahar, Shai Bergman, Mark Silberstein*

# Problems

- no I/O abstractions
- no address space management
- no page fault handling mechanisms
- not allow modifications to memory mappings for running GPU programs

# Motivation & Goals

## Build a software system:

To overcome the limitations of the limitation GPU's hardware virtual memory.

- Memory mapped files
- One can build an encrypted file system for GPUs by installing custom page fault handlers for encrypt- ing/decrypting file contents on-the-fly
- Shared memory system in a cluster of GPUs

The key element of our design is a new type of memory pointer we call *ActivePointer apointer*.

# Considerations & Challenges - From software

- **Main challenge:** Low translation overhead.
- **Efficient thread-level translation:** fast and convenient access, no deadlock when handling page fault.
- **Scalability.**

# Considerations & Challenges - From hardware

- **Coherence:** shared memory and L1 cache in each SM are not coherent, but stale mappings would result in page fault.
- **Asynchronous changes of page mappings.**

Constraining the ability of the paging system to revoke application access to a page at will is necessary to prevent asynchronous page mapping changes, eliminating the need for translation cache coherence.

# Design & Implementation - Principles

- **Active pages with fixed mapping.**

The page cannot be evicted from the page cache if it is being active used. Therefore each thread can cache the page mapping in its thread-private memory.

- **Keeping track of active pages.**

# Design & Implementation - Active Pointers (apointers)

```
int foo(){  
    //ptr initialized unlinked  
    APtr<float> ptr =  
        gvmmap(size, O_RDONLY, fd, foffset)  
    ptr += 10; // pointer arithmetics  
    float f1 = *ptr; // page fault on the first access  
    *ptr=25; // page fault free access via linked ptr  
}  
//ptr destroyed and unlinked
```

Figure 3: A simple example of using an *apointer* in GPU code. The *apointer* is initialized by calling the GPU version of `mmap()`, described in Section V.

# Design & Implementation - Active Pointers (apointers)

3 states: uninitialized, unlinked or linked

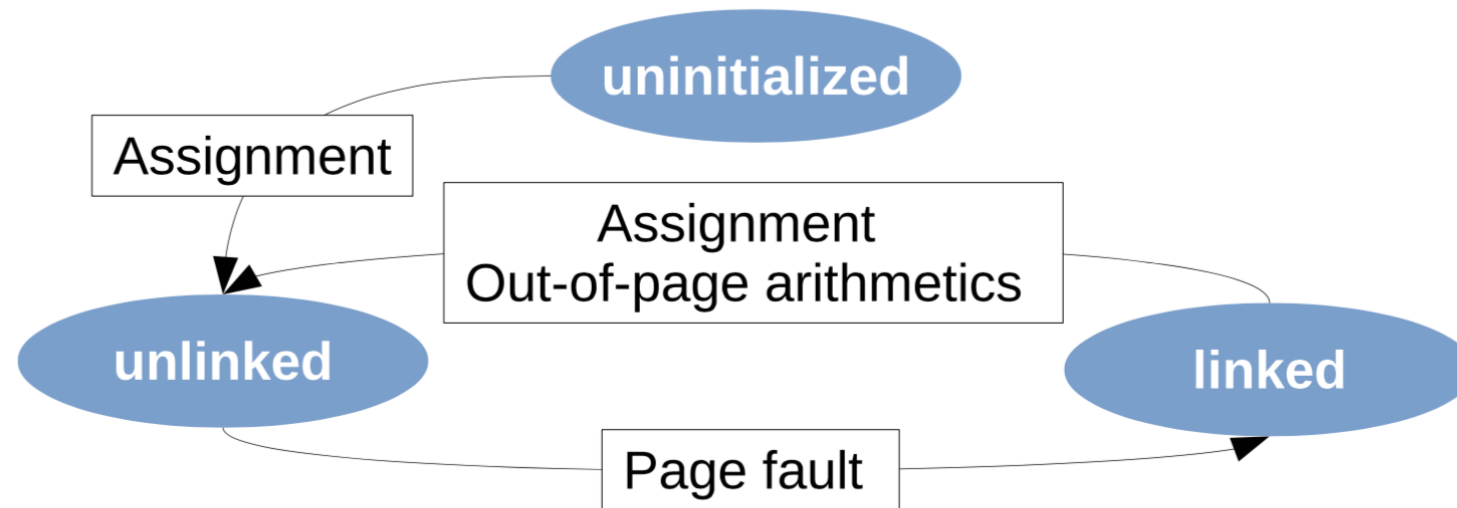


Figure 4: *Apointer* state transition diagram



# Design & Implementation - Active Pointers (apointers)

## Reference count

**Reference count (in the page cache):** the number of linked *apointers* holding the reference to that page.

- **Increment:** the page transitions to the linked state
- **Decrement:** it becomes unlinked or it is destroyed outside the program scope

# Design & Implementation - Active Pointers (apointers)

Data structure: translation field & metadata

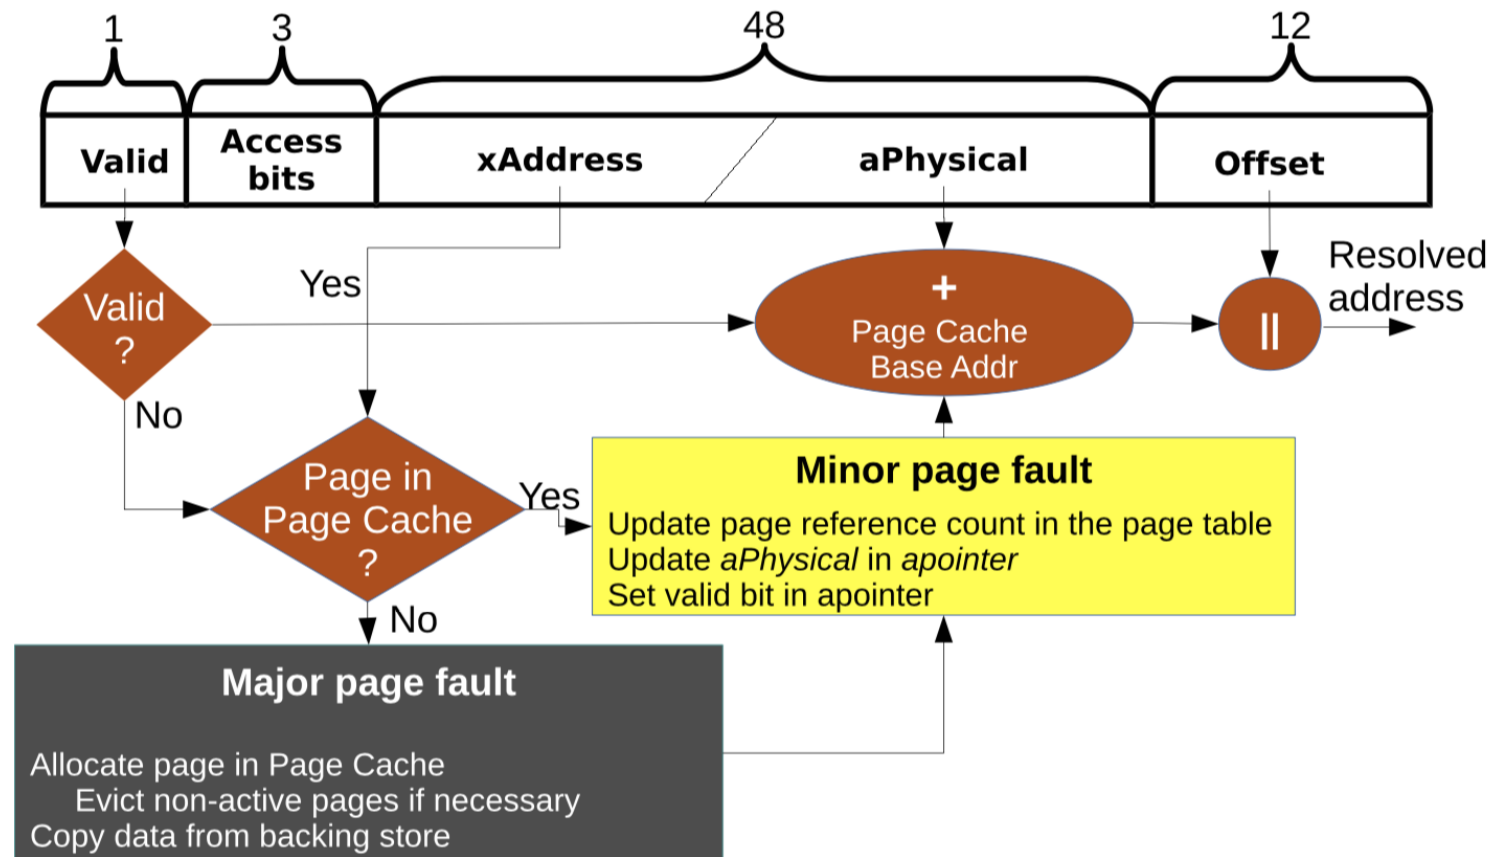


Figure 5: A functional diagram describing the use of *apointers*

# Design & Implementation - Thread-level address translation

- Design alternatives:
  - a *short apointer*: an *aphysical* address and `ndxAddress`
  - a *long apointer*: one of them
- Optimizing performance via speculative prefetch.
- Translation aggregation.

# Design & Implementation - Thread-level address translation

```
1 T dereference(aptr) {
2     // Test if there are page faults
3     isPagefault!=__all(aptr.valid);
4     if(isPagefault) pageFault(aptr);
5     return *(aptr.aPhys);
6 }
7 pageFault(aptr) {
8     // as long as there unhandled page faults
9     while(true) {
10        // Choose a leader in the group
11        warpLeader=__ffs(__ballot(!aptr.valid));
12        // No more pagefaults to handle
13        if(warpLeader == 0) break;
14        // Broadcast leader's backing store address to
15        // all threads
16        bAddr=__shfl(aptr.bAddr, warpLeader);
17
18        // Aggregate page reference count
19        isRequestHandling=(aptr.bAddr == bAddr);
20        pageRefCount=__popc(__ballot(requestHandling));
21        // Handle page fault using all threads in a warp
22        // Access locks only in warpLeader
23        // Update page reference count
24        aPhys =
25            HandlePageFault(warpLeader, bAddr, pageRefCount);
26        if(isRequestHandling) {
27            aptr.aPhys=aPhys;
28            aptr.valid=1;
29        }
30    }
31 }
```

Listing 1: Translation aggregation using CUDA warp primitives: `__all`: test if all warp threads satisfy a predicate. `__ballot`: fetch one bit across all warp threads, `__shfl`: send a word to all warp threads, `__ffs`: find the first set bit, and `__popc`: count the total number of set bits.

# Design & Implementation - Software TLB

- **Not necessary:** hardware registers hold the cache.
- Each threadblock maintains its own TLB for its threads.
- The TLB keeps the *threadblock-private* reference count for each cached page.

## Problems

- Count lost
- Deadlock
- Multiple TLBs ( global apointers )
- Overhead

# Evaluation & Analysis

- 2 X 6-core Intel i7-4960X CPUs at 3.6GHz, with 15MB L3/CPU, with power management and hyperthreading disabled for ensuring consistent results.
- A single GPU of the dual NVIDIA Tesla K80.
- Ubuntu Linux kernel 3.13.0-32 with CUDA SDK 7.0 and NVIDIA GPU driver 346.59.
- All baseline implementations require fewer than 64 registers/thread and do not spill registers.

## Evaluation & Analysis - Focus on

- Overhead of the software address translation layer on GPUs.
- The end-to-end performance of applications that use `ActivePointers` to map large datasets into GPU memory.

# Evaluation & Analysis - Apointer performance in page-fault free accesses

## Lantancy overhead

- 32 threads (one warp) where all the threads perform coalesced accesses to different offsets in one page.
- Each access involves a memory read and an increment operation.

Implementation	read	inc	read+ inc	read inc+rw
Raw access	225	32	257	257
Compiler	367 (+63%)	152 ( $\times 3.7$ )	519 (+101%)	585 (+127%)
Optimized PTX	282 (+25%)	–	434 (+69%)	544 (+111%)
Prefetching	271 (+20%)	–	423 (+65%)	435 (+75%)

Table I: GPU cycles when using *apointer* 4-byte read and increment (inc), separately and combined, and with page permission checks (rw), compared to the number of cycles when using a regular pointer (first row). Overhead is shown in parenthesis. Lower is better.

The most efficient apointer implementation uses 18 instructions vs. only 2 for a simple pointer increment.



# Evaluation & Analysis - Apointer performance in page-fault free accesses

## Throughput overhead

- Run hundreds of warps to **saturate** all compute units in the mapped files and differs from the standard file system access pattern.
- Memory tiling

Implementation	4-byte	4-byte+rw	8-byte
Compiler	99.7GB/s (65.4%)	97.7 (64.1%)	148.7 (97.6%)

Table II: Memory bandwidth in GB/s achieved by memory copy kernel, compared to the maximum achievable bandwidth of 152GB/s (in parentheses). Higher is better.


# Evaluation & Analysis - Apointer performance in page-fault free accesses

## Free- computation bubble


- NVIDIA K80 GPU issues  $2056 \times 10^9$  instructions per second per GPU
- Memory bandwidth:  $240 \times 10^9$  bytes/sec
- Free-computation bubble: 8.6 instructions per byte of memory traffic

# Evaluation & Analysis - Compute-intensive workloads

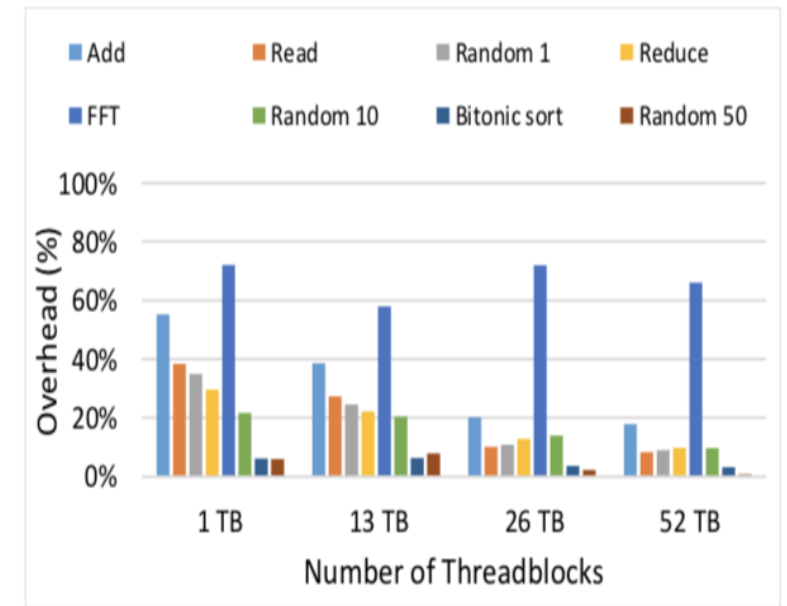
## 4-byte reads



(a) 4-byte reads



(b) 16-byte reads



(c) 4-byte reads with GPUfs

Figure 6: *Apointer* overheads as a function of GPU occupancy (in number of threadblocks) (a) for 4-byte reads and (b) 16-byte reads excluding GPUfs (§ VI-B), and (c) for 4-byte reads with GPUfs (§ VI-C). Lower is better.

# Evaluation & Analysis - Page cache and apointers

## Page fault

Implemetation	Minor Pagefault	Major Pagefault
<i>Apointer</i> Short	20%	No observable overhead
<i>Apointer</i> Long	24%	No observable overhead
no TLB	13%	No observable overhead

Table III: The overhead of short *apointer*, long *apointers*, and long *apointers* without TLB, with major and minor page faults. Lower is better.

The best performance, however, is achieved without the TLB with long *apointer* because it avoids the overheads of TLB updates.

# Evaluation & Analysis - Page cache and apointers

## Effects of TLB size





Figure 7: Read access times in cycles per page, as a function of unique pages accessed per threadblock, for different TLB sizes. Lower is better.


# Evaluation & Analysis - Compute-intensive workloads with page cache



(c) 4-byte reads with GPUfs

# Evaluation & Analysis - End-to-end application performance

The use of *apointers* does not introduce any measurable overheads over the fastest GPUfs-only implementation, and therefore achieves both high end-to-end performance and programming simplicity in this complex I/O-intensive application.



# Discussion

- Register pressure
- Compiler support
- Instructions for boundary checking and pointer increment
- I/O preemption



# The Last Word

- Page fault is available since Pascal
- DMA